

# Deep Learning with Microsoft Azure:

Margaryta Ostapchuk  
Nnanna Oriek  
Sabrina Smai

# Table of Contents

## Machine Learning Requirements

Applications, Characteristics and Requirements

End-to-End lifecycle and processes

Deep Learning: Additional Requirements

## Azure Machine Learning Service

Artifacts: Workspace, Experiments, Compute, Models, Images, Deployment and Datastore

Concepts: Model Management, Pipelines

## Azure Automated Machine Learning

## Distributed Training with Azure ML Compute

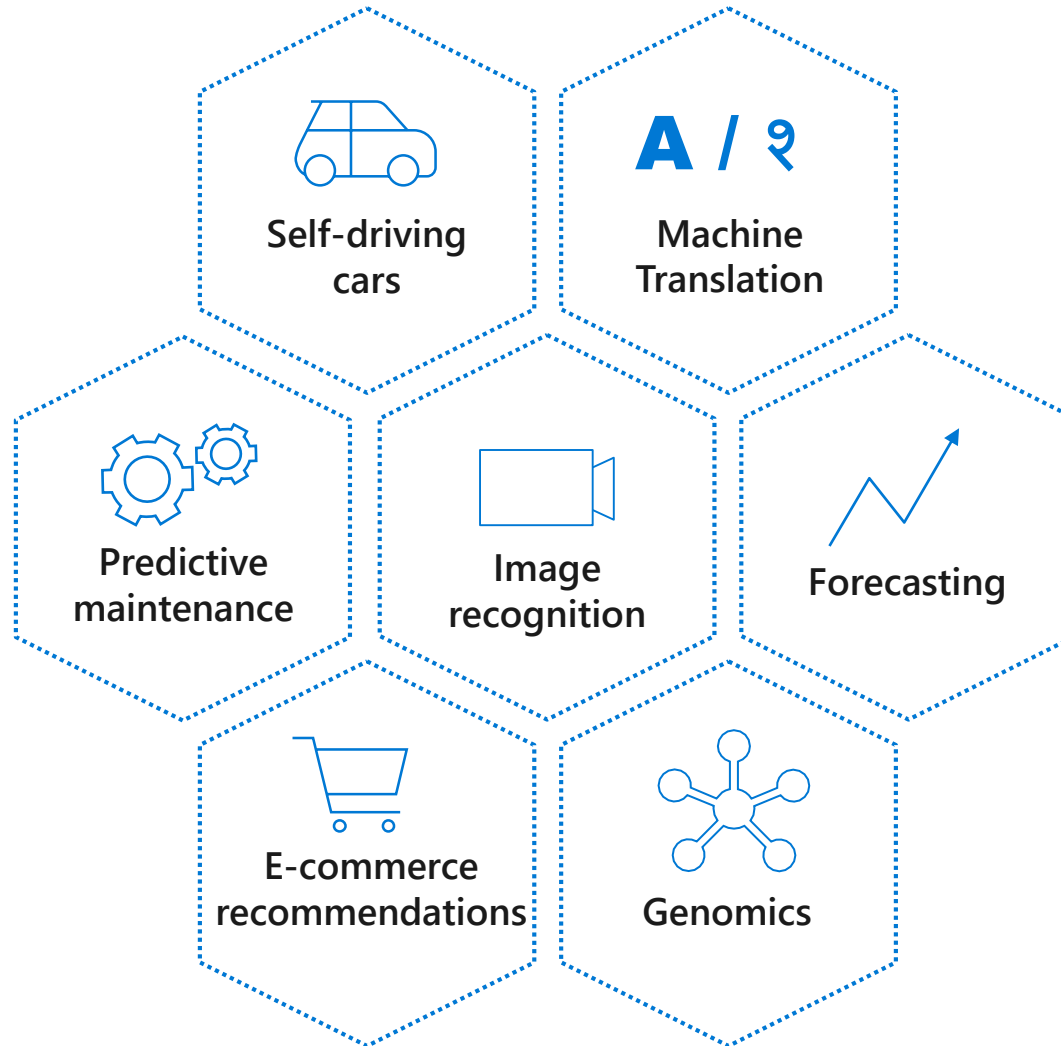
## Deploy Azure Models

To Edge Devices

## Azure Machine Learning Demo:

# Machine Learning

## Applications

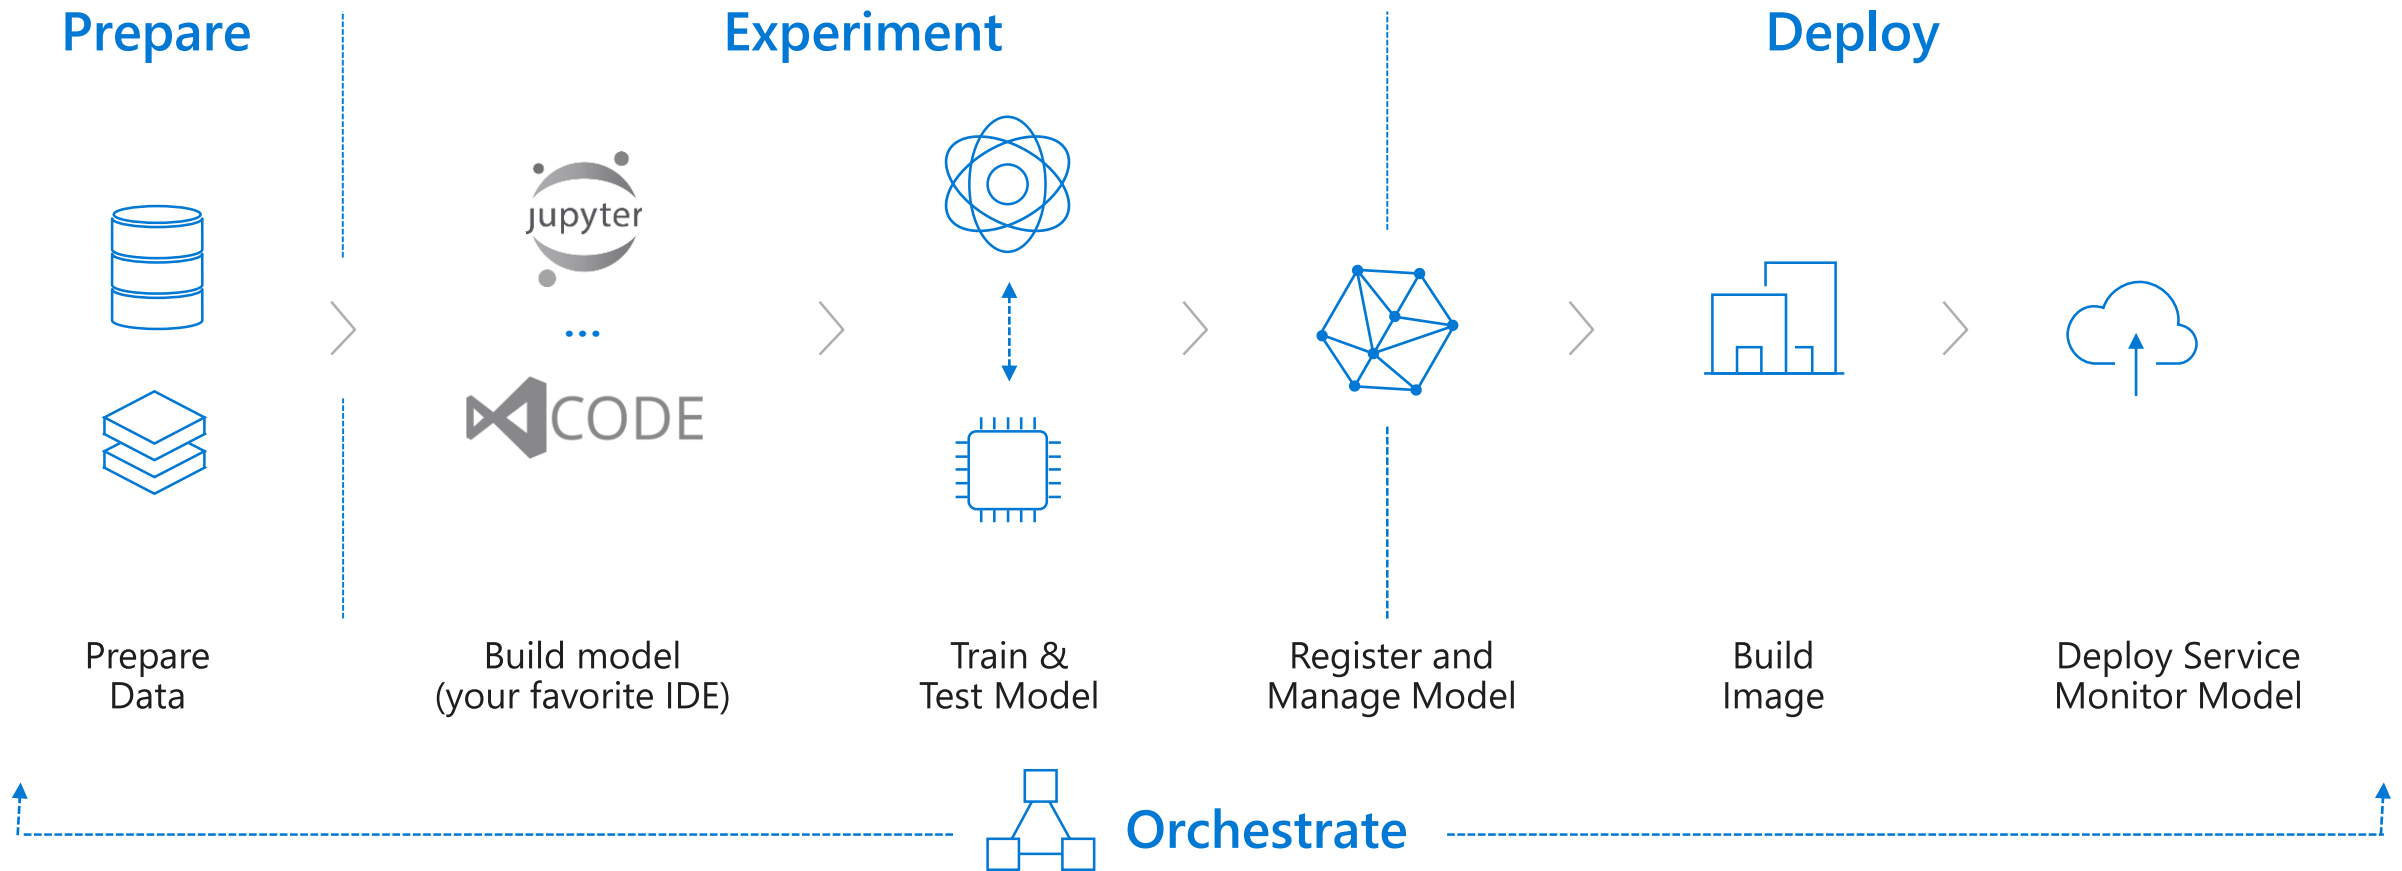





# Requirements of an advanced ML Platform

# Machine Learning

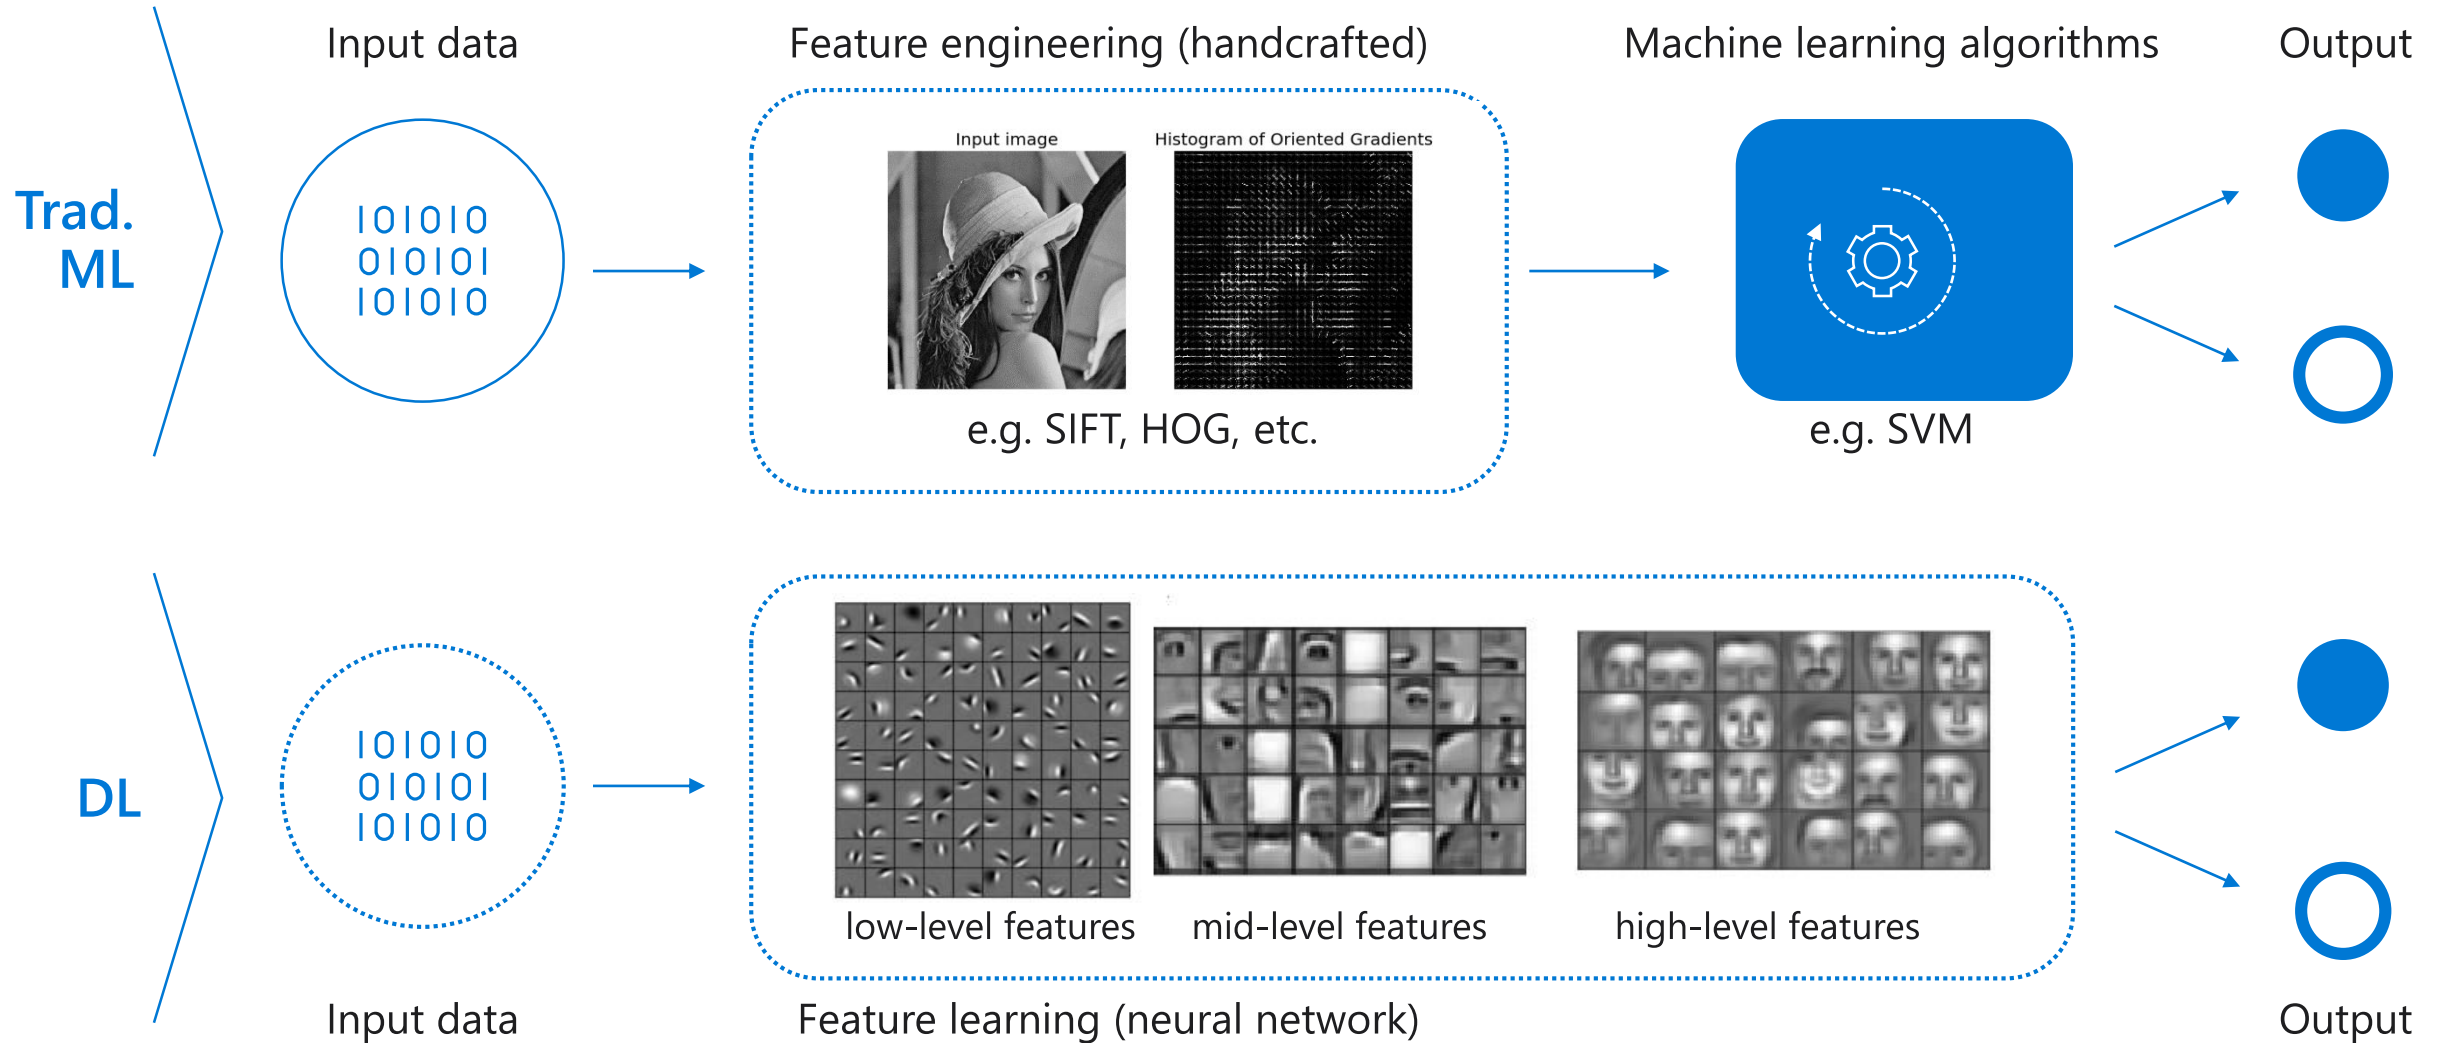
Typical E2E Process






Deep Learning places  
additional requirements

# Traditional ML versus DL



# Characteristics of Deep Learning



101010  
010101  
101010

Massive amounts of training data



Excels with raw, unstructured data



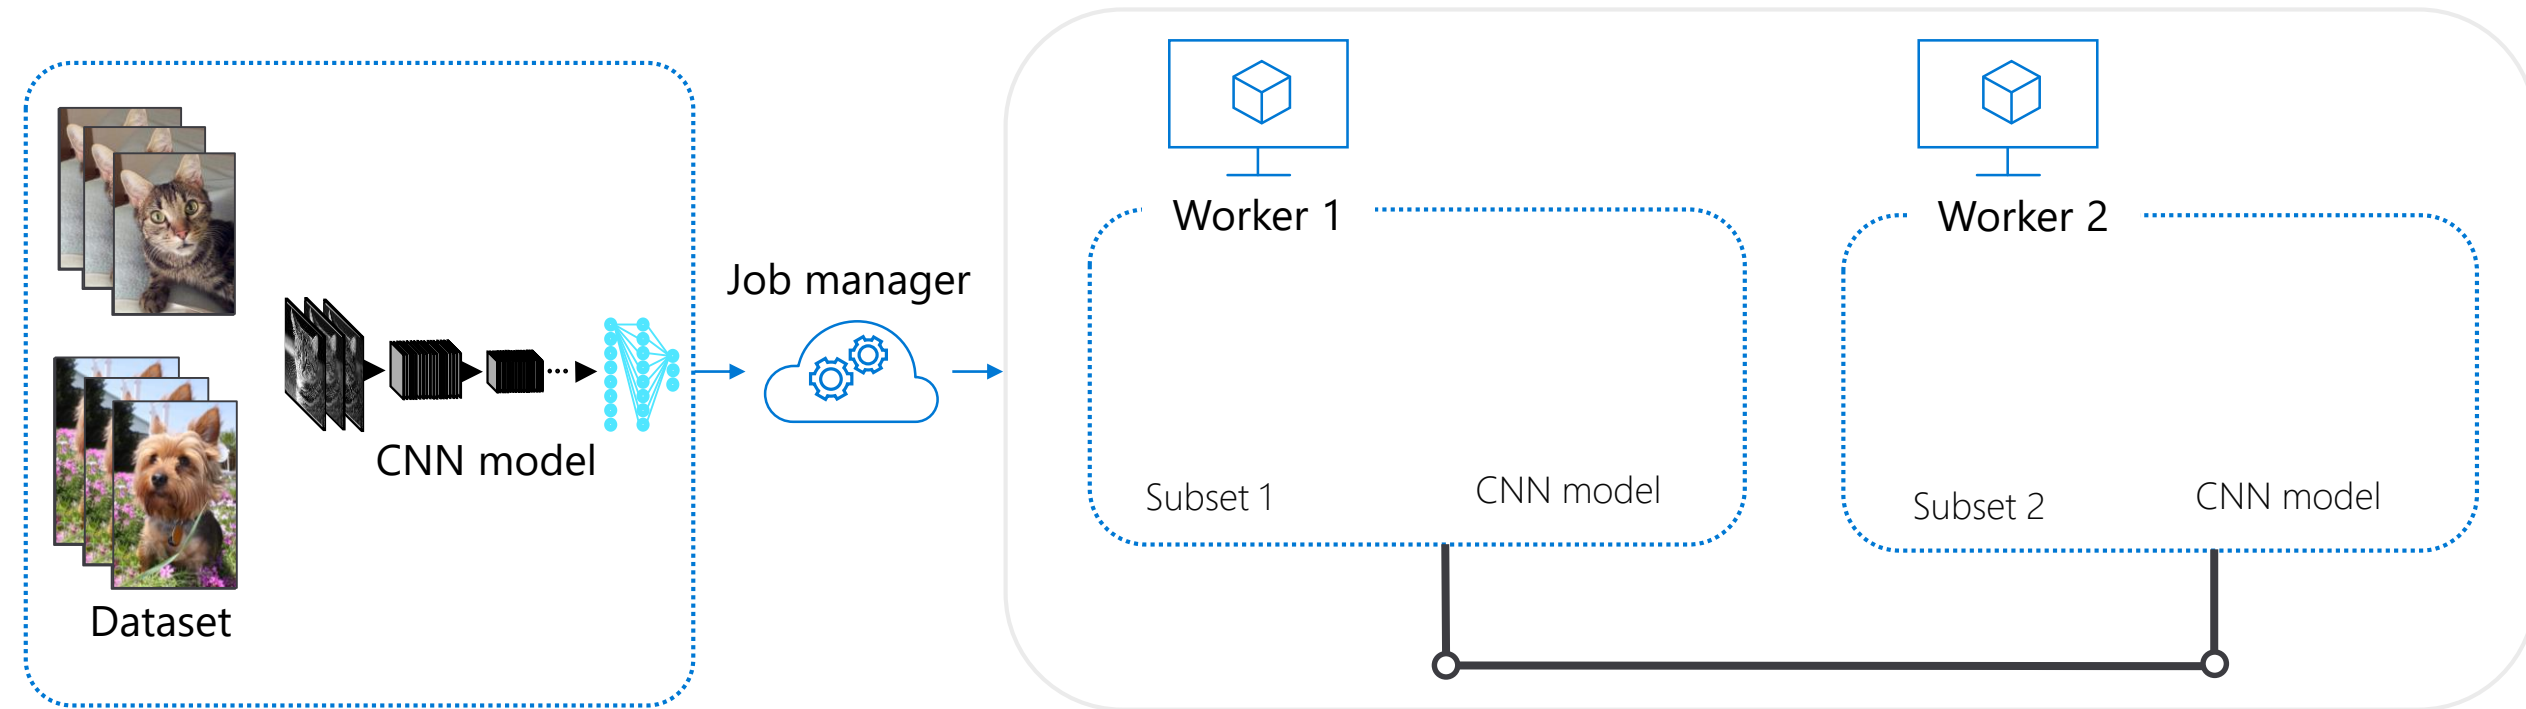
Automatic feature extraction



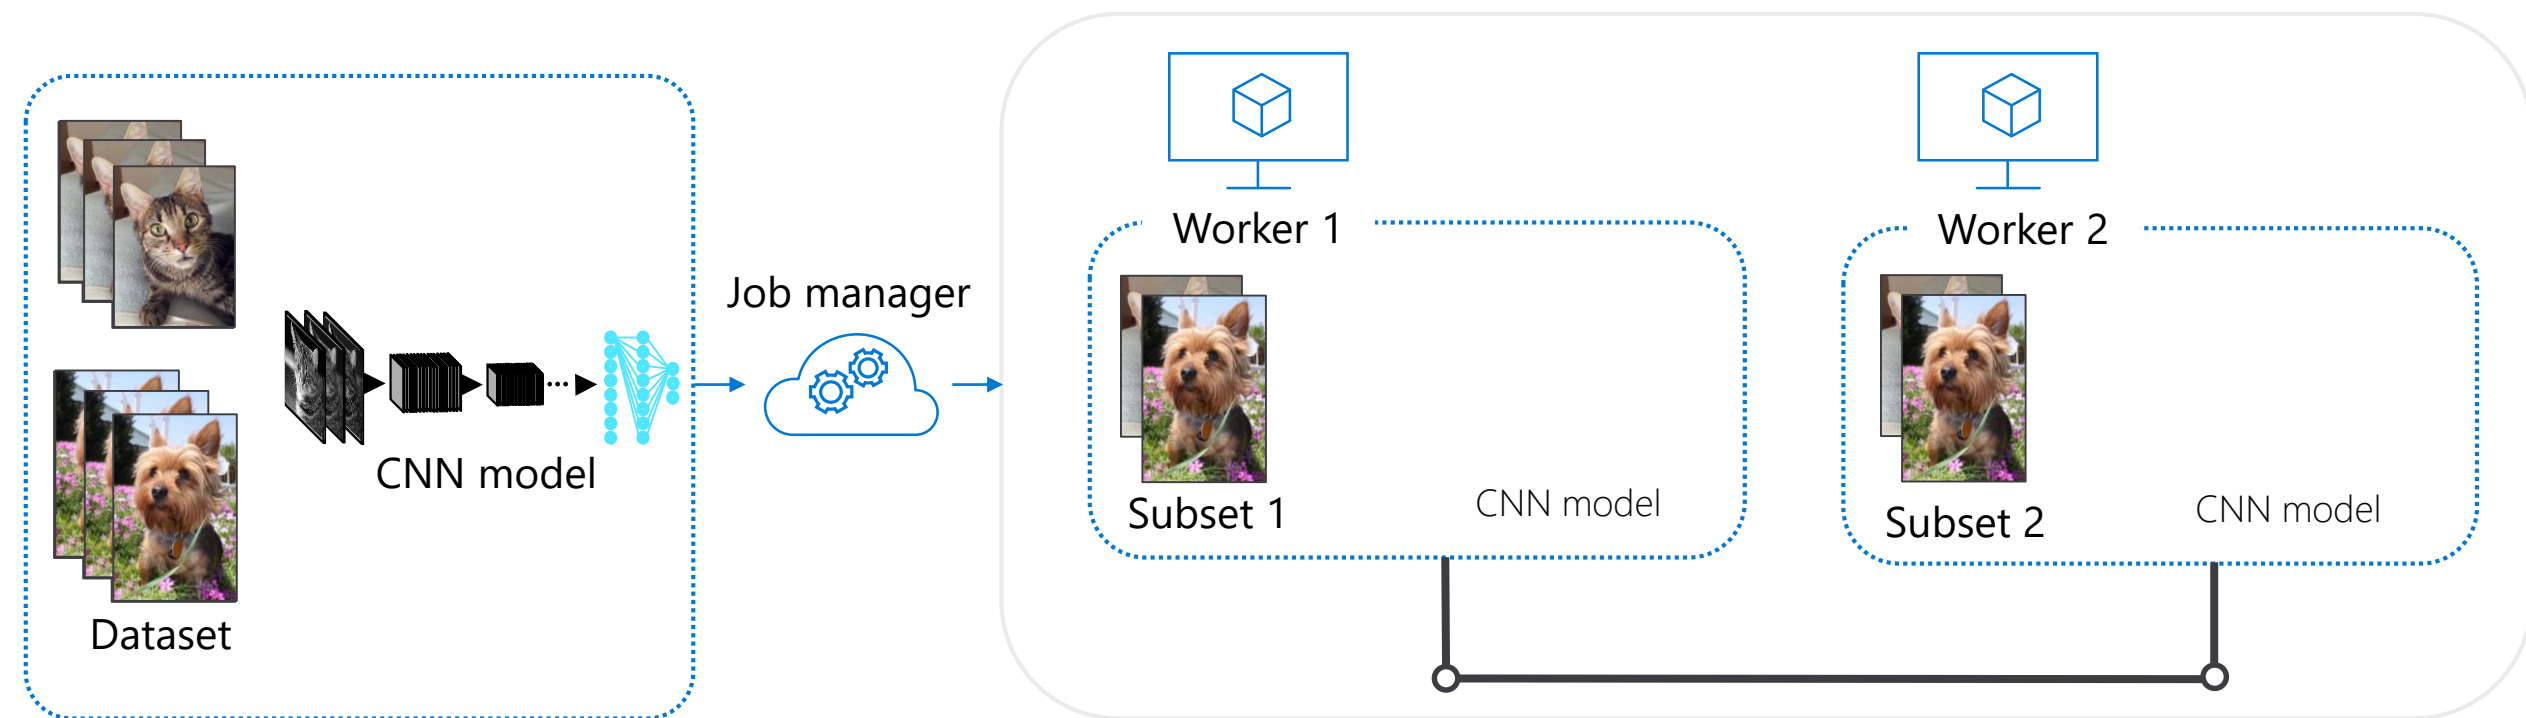
Computationally expensive



# Distributed training mode: Data parallelism



# Distributed training mode: Model parallelism



# Required Deep Learning Frameworks

## Popular Deep Learning Frameworks



TensorFlow



PyTorch



Scikit-Learn



MXNet



Chainer




Keras




ONNX

Community project created by Facebook and Microsoft  
Use the best tool for the job. Train in one framework  
and transfer to another for inference



# Deep Learning

## Three Additional Requirements



1. Distributed Training on multi-node clusters
2. Support for advanced processors: TPUs GPUs FPGAs
3. Support for Deep Learning Frameworks:



**Azure offers a comprehensive  
AI/ML platform that meets—and  
exceeds—requirements**

# Machine Learning on Azure

## Domain specific pretrained models

To reduce time to market



Vision



Speech



Language



Search

## Familiar Data Science tools

To simplify model development



PyCharm



Jupyter



Visual Studio Code



Command line

## Popular frameworks

To build advanced deep learning solutions



Pytorch



TensorFlow



Scikit-Learn



Onnx

## Productive services

To empower data science and development teams



Azure  
Databricks



Azure Machine  
Learning



Machine  
Learning VMs

## Powerful infrastructure

To accelerate deep learning



CPU



GPU



FPGA



From the Intelligent Cloud to the Intelligent Edge



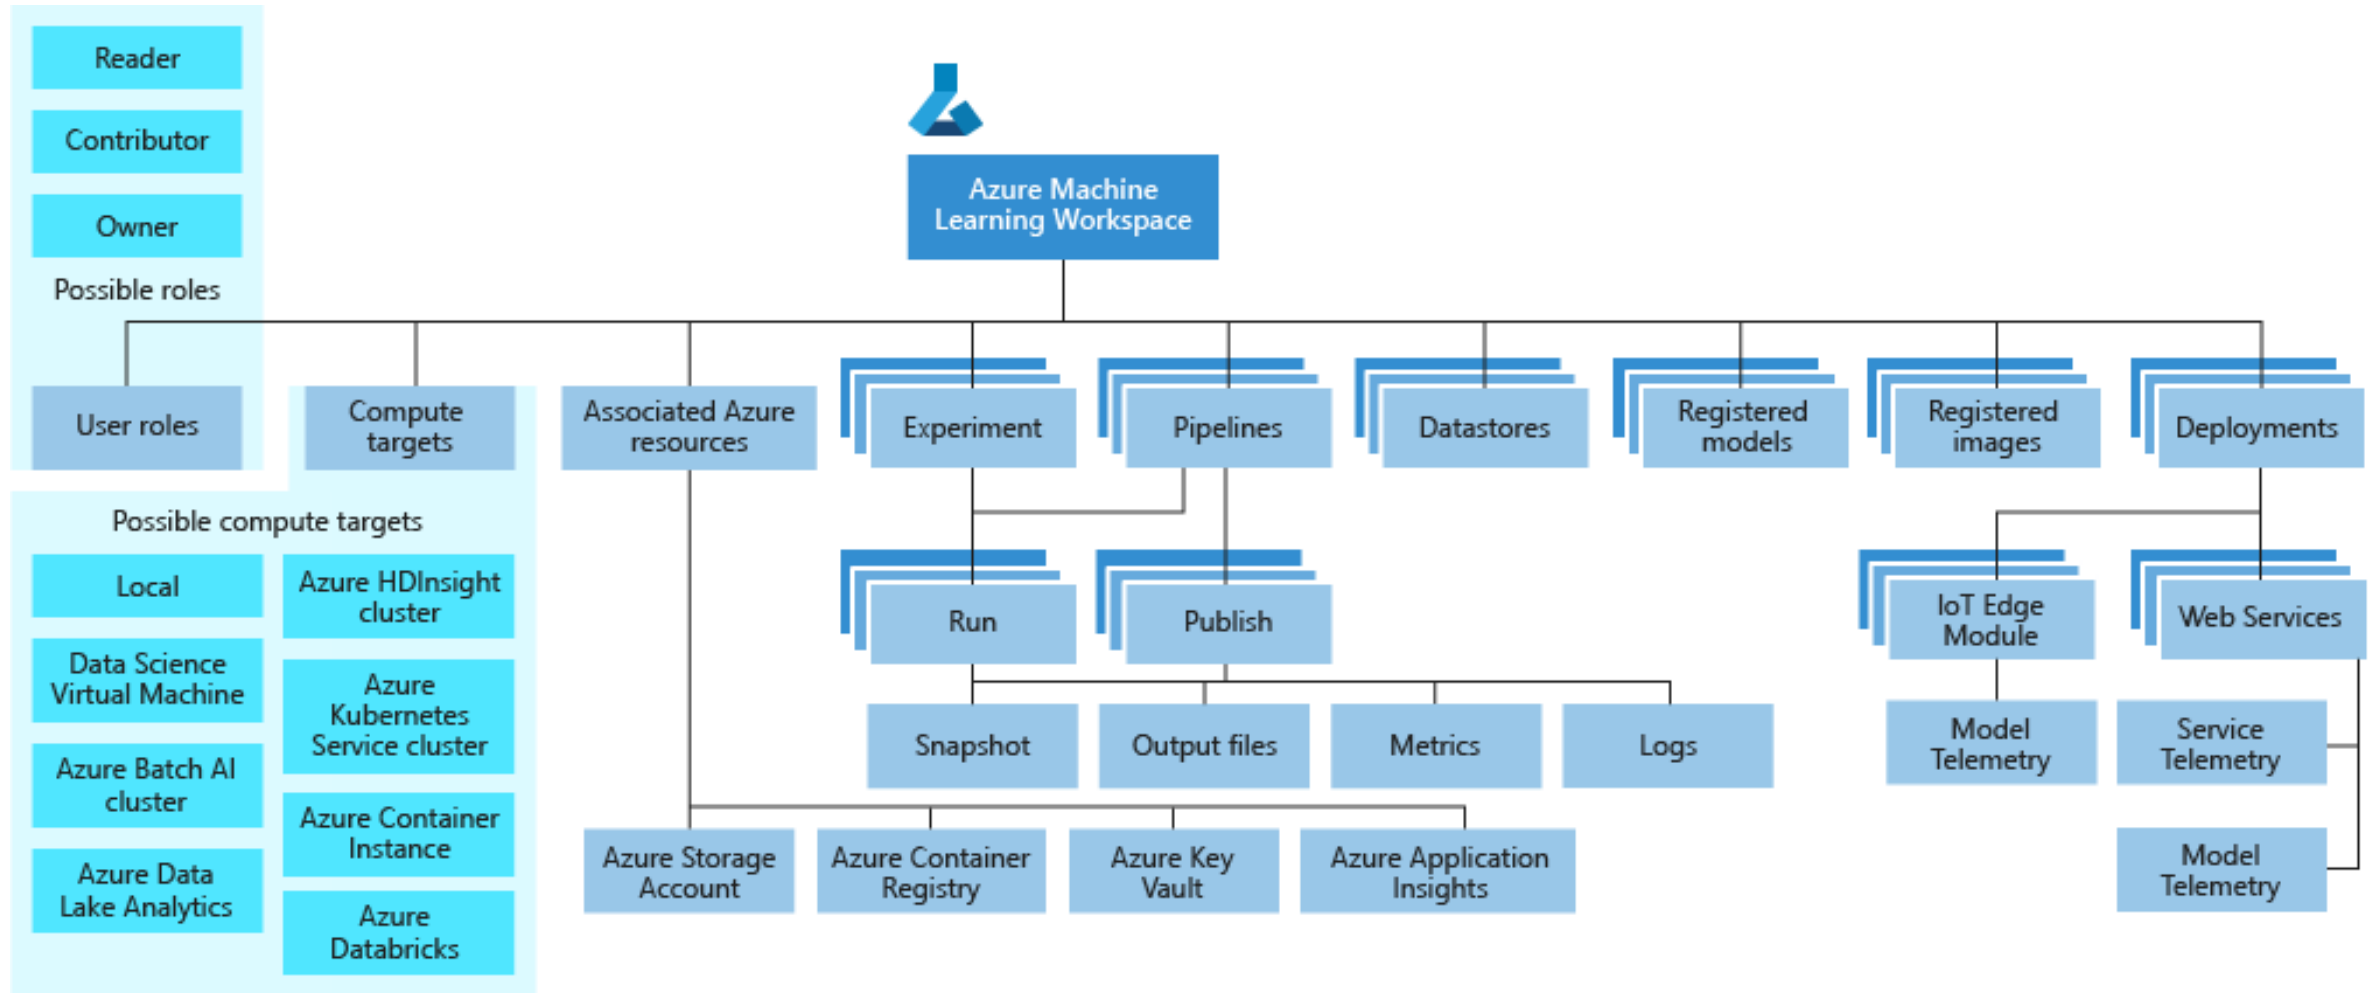
# What is Azure Machine Learning service?



That enables  
you to:

- ✓ Prepare Data
- ✓ Build Models
- ✓ Train Models
- ✓ Manage Models
- ✓ Track Experiments
- ✓ Deploy Models

# Azure ML service Workspace Taxonomy





# Azure ML

## Steps

1. Snapshot folder and send to experiment

6. Stream stdout, logs, metrics

Experiment

6. Stream stdout, logs, metrics

7. Copy over outputs

2. Create docker image

Docker Image

5. Launch script


Compute Target

3. Deploy docker and snapshot to compute

4. Mount datastore to compute

Data Store

Azure ML  
Workspace



My Computer

## Step 1 – Create a workspace

```
from azureml.core import Workspace
ws = Workspace.create(name='myworkspace',
                     subscription_id='<azure-subscription-id>',
                     resource_group='myresourcegroup',
                     create_resource_group=True,
                     location='eastus2' # or other supported Azure region
)

# see workspace details
ws.get_details()
```

## Step 2 – Create an Experiment

Create an experiment to track the runs in the workspace. A workspace can have multiple experiments

```
experiment_name = 'my-experiment-1'

from azureml.core import Experiment
exp = Experiment(workspace=ws, name=experiment_name)
```

## Step 3 – Create remote compute target

```
# choose a name for your cluster, specify min and max nodes
compute_name = os.environ.get("BATCHAI_CLUSTER_NAME", "cpucluster")
compute_min_nodes = os.environ.get("BATCHAI_CLUSTER_MIN_NODES", 0)
compute_max_nodes = os.environ.get("BATCHAI_CLUSTER_MAX_NODES", 4)

# This example uses CPU VM. For using GPU VM, set SKU to STANDARD_NC6
vm_size = os.environ.get("BATCHAI_CLUSTER_SKU", "STANDARD_D2_V2")

provisioning_config = AmlCompute.provisioning_configuration(
    vm_size = vm_size,
    min_nodes = compute_min_nodes,
    max_nodes = compute_max_nodes)

# create the cluster
print(' creating a new compute target... ')
compute_target = ComputeTarget.create(ws, compute_name, provisioning_config)

# You can poll for a minimum number of nodes and for a specific timeout.
# if no min node count is provided it will use the scale settings for the cluster
compute_target.wait_for_completion(show_output=True,
                                   min_node_count=None, timeout_in_minutes=20)
```

Zero is the default.  
If min is zero then  
the cluster is  
automatically  
deleted when no  
jobs are running  
on it.

## Step 4 – Upload data to the cloud

First load the compressed files into numpy arrays. Note the '`load_data`' is a custom function that simply parses the compressed files into numpy arrays.

```
# note that while loading, we are shrinking the intensity values (X) from 0-255 to 0-1 so that the
model converge faster.
X_train = load_data('./data/train-images.gz', False) / 255.0
y_train = load_data('./data/train-labels.gz', True).reshape(-1)

X_test = load_data('./data/test-images.gz', False) / 255.0
y_test = load_data('./data/test-labels.gz', True).reshape(-1)
```

Now make the data accessible remotely by uploading that data from your local machine into Azure so it can be accessed for remote training. The files are uploaded into a directory named `mnist` at the root of the datastore.

```
ds = ws.get_default_datastore()
print(ds.datastore_type, ds.account_name, ds.container_name)

ds.upload(src_dir='./data', target_path='mnist', overwrite=True, show_progress=True)
```

We now have everything you need to start training a model.

## Step 5 – Train a local model

Train a simple logistic regression model using scikit-learn locally. This should take a minute or two.

```
%%time from sklearn.linear_model import LogisticRegression
clf = LogisticRegression()
clf.fit(X_train, y_train)

# Next, make predictions using the test set and calculate the accuracy
y_hat = clf.predict(X_test)
print(np.average(y_hat == y_test))
```

You should see the local model accuracy displayed. [It should be a number like 0.915]

## Step 6 – Train model on remote cluster

To submit a training job to a remote you have to perform the following tasks:

- 6.1: Create a directory
- 6.2: Create a training script
- 6.3: Create an estimator object
- 6.4: Submit the job

### Step 6.1 – Create a directory

Create a directory to deliver the required code from your computer to the remote resource.

```
import os
script_folder = './sklearn-mnist' os.makedirs(script_folder, exist_ok=True)
```

## Step 6.2 – Create a Training Script (1/2)

```
%writefile $script_folder/train.py
# load train and test set into numpy arrays
# Note: we scale the pixel intensity values to 0-1 (by dividing it with 255.0) so the model can
# converge faster.
# 'data_folder' variable holds the location of the data files (from datastore)
Reg = 0.8 # regularization rate of the logistic regression model.
X_train = load_data(os.path.join(data_folder, 'train-images.gz'), False) / 255.0
X_test  = load_data(os.path.join(data_folder, 'test-images.gz'), False) / 255.0
y_train = load_data(os.path.join(data_folder, 'train-labels.gz'), True).reshape(-1)
y_test  = load_data(os.path.join(data_folder, 'test-labels.gz'), True).reshape(-1)
print(X_train.shape, y_train.shape, X_test.shape, y_test.shape, sep = '\n')

# get hold of the current run
run = Run.get_context()
#Train a logistic regression model with regularizaion rate of 'reg'
clf = LogisticRegression(C=1.0/reg, random_state=42)
clf.fit(X_train, y_train)
```

## Step 6.2 – Create a Training Script (2/2)

```
print('Predict the test set')
y_hat = clf.predict(X_test)

# calculate accuracy on the prediction
acc = np.average(y_hat == y_test)
print('Accuracy is', acc)

run.log('regularization rate', np.float(args.reg))
run.log('accuracy', np.float(acc)) os.makedirs('outputs', exist_ok=True)

# The training script saves the model into a directory named 'outputs'. Note files saved in the
# outputs folder are automatically uploaded into experiment record. Anything written in this
# directory is automatically uploaded into the workspace.
joblib.dump(value=clf, filename='outputs/sklearn_mnist_model.pkl')
```



## Step 6.3 – Create an Estimator

An estimator object is used to submit the run.

```
from azureml.train.estimator import Estimator

script_params = { '--data-folder': ds.as_mount(), '--regularization': 0.8 }

est = Estimator(source_directory=script_folder,
                script_params=script_params,
                compute_target=compute_target,
                entry_script='train.py',
                conda_packages=['scikit-learn'])
```

The directory that contains the scripts. All the files in this directory are uploaded into the cluster nodes for execution

Name of  
estimator

Python Packages  
needed for training

Training Script  
Name


Compute  
target (Batch AI  
in this case)

Parameters required  
from the training script

## Step 6.4 – Submit the job to the cluster for training

```
run = exp.submit(config=est)
run
```

# What happens after you submit the job?



## Image creation

A Docker image is created matching the Python environment specified by the estimator. The image is uploaded to the workspace. Image creation and uploading takes about 5 minutes.

This happens once for each Python environment since the container is cached for subsequent runs. During image creation, logs are streamed to the run history. You can monitor the image creation progress using these logs.



## Scaling

If the remote cluster requires more nodes to execute the run than currently available, additional nodes are added automatically. Scaling typically takes about 5 minutes.



## Running

In this stage, the necessary scripts and files are sent to the compute target, then data stores are mounted/copied, then the entry\_script is run. While the job is running, stdout and the ./logs directory are streamed to the run history. You can monitor the run's progress using these logs.



## Post-Processing

The ./outputs directory of the run is copied over to the run history in your workspace so you can access these results.

# Azure ML Concept

## Model Management


Model Management in Azure ML usually involves these four steps

**Step 1:** Register Model using the Model Registry

**Step 2:** Register Image using the Image Registry (the Azure Container Registry)

**Step 3:** Deploy the Image to cloud or to edge devices

**Step 4:** Monitor models—you can monitor input, output, and other relevant data from your model.



# Azure ML Artifact

## Deployment

Deployment is an instantiation of an image. Two options:



A dashed blue line originates from the text 'Deployment is an instantiation of an image. Two options:'. It extends horizontally to the right and then splits into two vertical dashed lines, each ending in a downward-pointing arrowhead. The left arrow points to the 'Web service' header, and the right arrow points to the 'IoT Module' header.

### Web service

A deployed web service can run on Azure Container Instances, Azure Kubernetes Service, or field-programmable gate arrays (FPGA).

Can receive scoring requests via an exposed a load-balanced, HTTP endpoint.

Can be monitored by collecting Application Insight telemetry and/or model telemetry.

Azure can automatically scale deployments.

### IoT Module

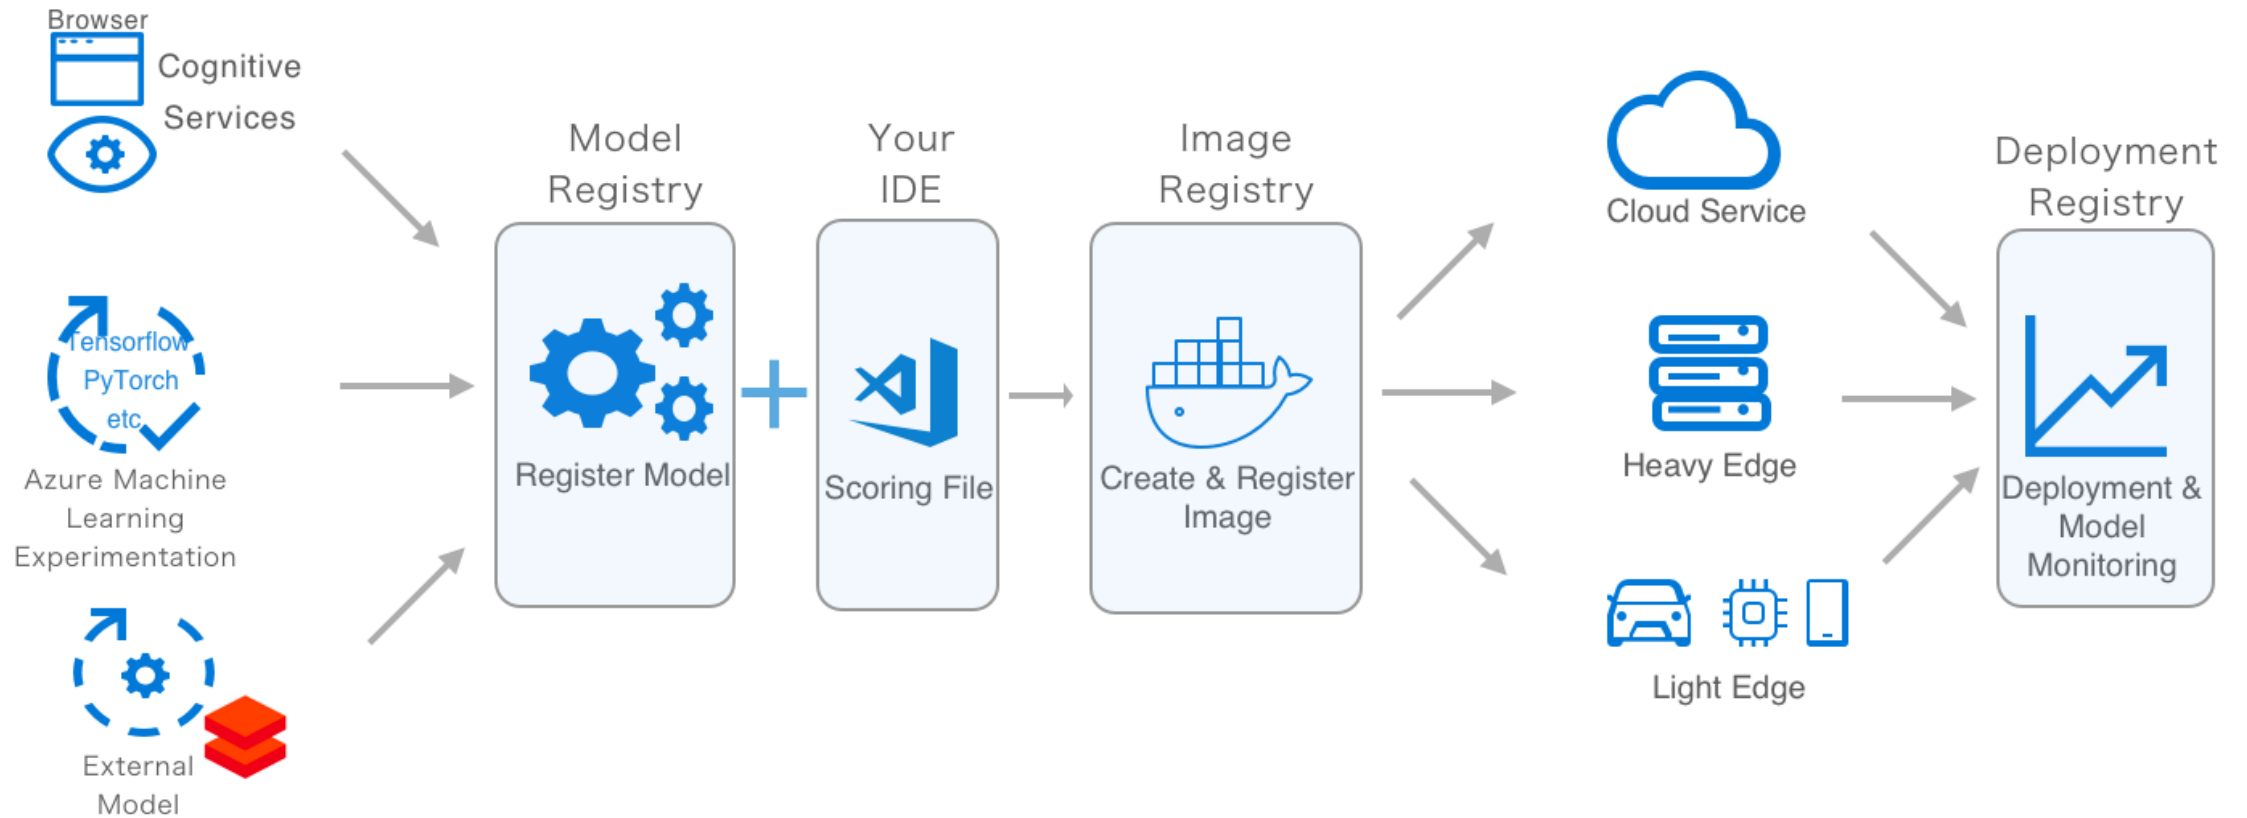
A deployed IoT Module is a Docker container that includes the model, associated script and additional dependencies.

Is deployed using **Azure IoT Edge** on edge devices.

Can be monitored by collecting Application Insight telemetry and/or model telemetry.


Azure IoT Edge will ensure that your module is running and monitor the device that is hosting it.

# Azure ML: How to deploy models at scale

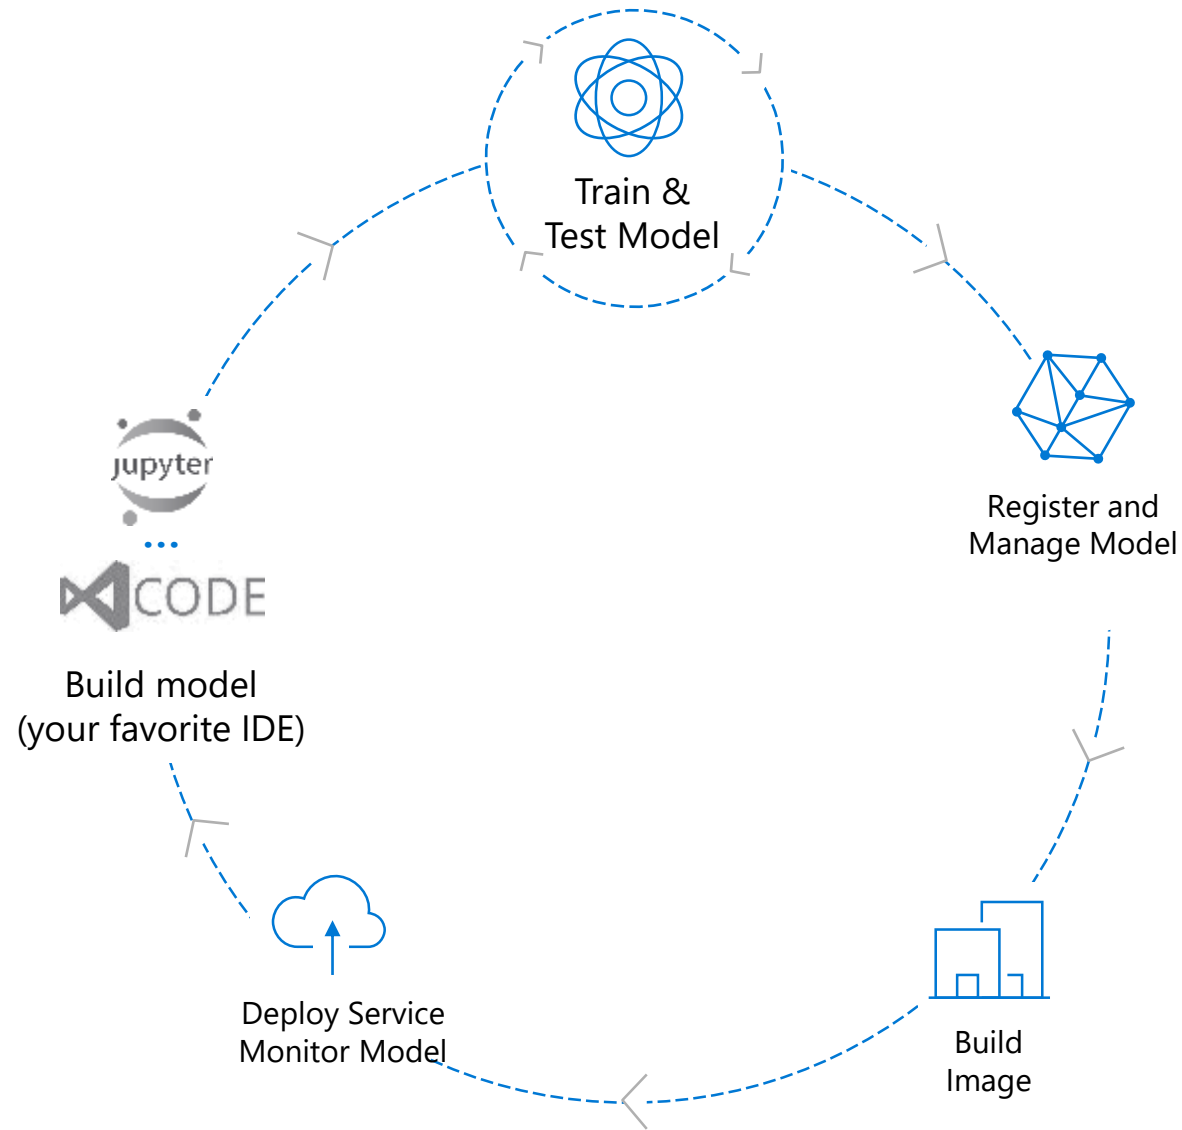


# DevOps loop for data science: AML + DevOps

Prepare



Prepare  
Data






**Azure Automated Machine Learning  
'simplifies' the creation and selection  
of the optimal model**

# Automated ML

## Current Capabilities

Category		Value
ML Problem Spaces		Classification Regression Forecasting
Frameworks		Scikit Learn
Languages		Python
Data Type and Data Formats		Numerical Text Scikit-learn supported data formats (Numpy, Pandas)
Data sources		Local Files, Azure Blob Storage
<a href="#">Compute Target</a>	Automated Hyperparameter Tuning	Azure ML Compute (Batch AI), Azure Databricks
	Automated Model Selection	Local Compute, Azure ML Compute (Batch AI), Azure Databricks





# Distributed Training with Azure ML Compute

# Distributed Training with Azure ML Compute

You submit a model training 'job' – the infrastructure is managed for you.

Jobs run on a native VM or Docker container.

Supports Low priority (Cheaper) or Dedicated (Reliable) VMS.

Auto-scales: Just specify min and max number of nodes.


*If min is set to zero, cluster is deleted when no jobs are running; so pay only for job duration.*

Works with most popular frameworks and multiple languages.

Supports [distributed training with Horovod](#).

Cluster can be shared; multiple experiments can be run in parallel.

Supports most VM Families, including latest NVidia GPUs for DL model training.

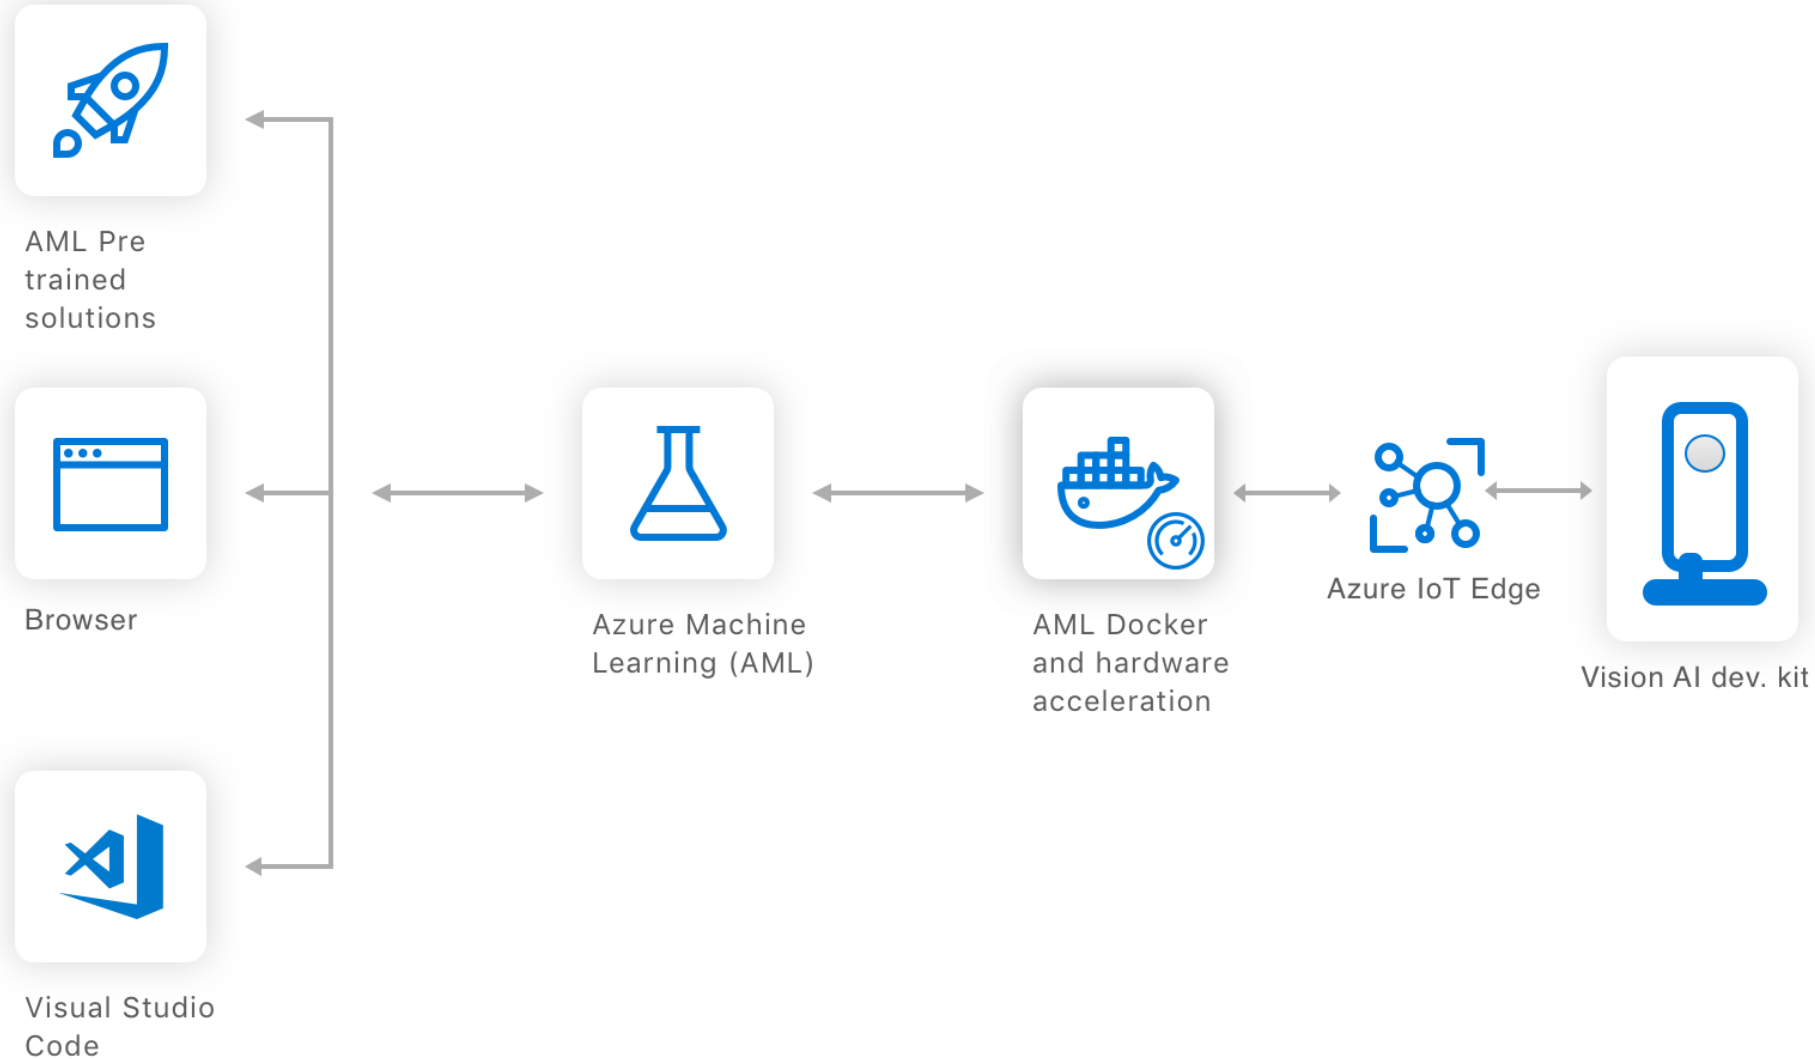




**Deploy to IoT Edge**

# Vision AI Development kit

## System Architecture



# What Azure Machine Learning Service

- End-to-End ML Life-cycle Management
- Model Interpretability-- InterpretML
- Fairness -- FairLearn
- DataDrift
- MLOPS – Integration with Azure Devops
- Build Anywhere, Manage and Deploy with AML
- Access to Compute on Demand
- Manage cost with Pipeline Reuse
- Low Priority VM
- Full Supports OSS and Enterprise



Demo